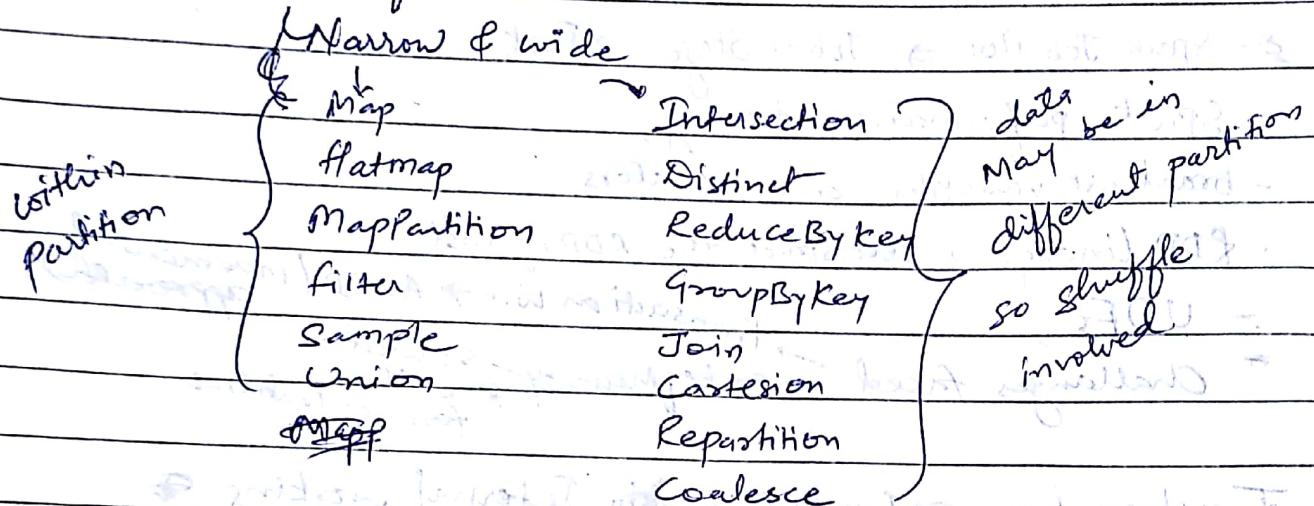


# Spark

## Spark Core

- ① Definition ② advantages over hadoop; Architecture
- ✓ ③ features of spark ④ ecosystem components
- I2 P3 C1 ⑤ RDD meaning, ways to create, features of RDD
- ⑥ RDD Transformations



## ⑦ Actions -

- first
- take
- reduce
- count
- collect
- cache → action
- saveAsTextfile
- ⑧ Limitation of RDD
  - ↳ no optimization - no structure
  - garbage collection issue,
  - performance degrade with bigger RDD.
  - No static typing

## ⑨ RDD Persistence & caching - benefits & modes & unpersist

↳ MO, MOS, MAD, MADS, DO, MO2 & MAD2, OFF-HEAP

## ⑩ Overcome the limitations of RDD.

↳ serialized

① Use dataset & dataframes (catalyst optimizer & Tungsten Exe. Engine)

② compile type check using DS & DF

③ DS & DF → more control over data using structure

## ⑪ Transformations - map, flatmap, mapPartitions, mapPartitionWithIndex,

Schema filter, union, intersection, distinct, reduceByKey, groupByKey, sortByKey  
must be same join, coalesce, cogroup, cartesian, pipe, repartition, combineByKey  
pair RDD only repartitionAndSortWithinPartitions, mapValues, flatMapValues,  
process RDD with shell script

$\nearrow$  unordered  $\rightarrow$  ordered, head

Spark Actions - count, collect, take, top, countByValue, reduce, fold, aggregate, foreach, takeSample, takeOrdered, saveAsTextfile, SaveAsSequencefile, saveAsObjectfile, countByKey, collectAsList, foreachPartition, co-partitioned API

co-partitioned API

API

→ Spark Job Flow  $\Rightarrow$  Job  $\rightarrow$  Stage  $\rightarrow$  Task

- Spark performance tuning,  
- broadcast variables, accumulators

- RDD lineage  $\rightarrow$  reconstruct the RDD if lost

- UDFs,  $\rightarrow$  Transaction live  $\rightarrow$  ACID (incremental approach)  
→ Challenges faced  $\leftarrow$  outofmemory solved  
with foreach partition

Transformations on fairRDD:-  $\rightarrow$  Internal working

② ReduceByKey, foldByKey, combineByKey, partitionBy (new HashPartitioner(2)), AggregateByKey.

Actions on PairRDD :- CountByKey, CollectAsMap, lookup(key)

Spark Partitioning -

→ Input scheme  $\rightarrow$  assigned to address of partitions  $\rightarrow$  RDD

→ RDD to partitions with mapping

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

→ Input file (mining log) converted to tuples sets

Spark - spark architecture, RDD, transformation & actions, advantages, features, limitation of RDD, persistence Narrow wide

Spark

SQL - Spark dataframes, datasets, difference, Optimization engine, Java API, spark Shuffling Data partitioning, Row format, broadcasts, accumulators

Spark Configuration properties, compression technique

Spark Streaming -

Challanges -

→ OOM → for Each Partition

Map Each Partition with Index → collect.

Hive → comma separated values → used Lateral view & explode

Code redundancy → modularised codes common utility

Spark tuning → persist RDD, remove collects

Used Accumulators V2, calculation

Latest data → Ranking windowing

of executor memory

~~10.97-101. PG: 9~~, ~~reading~~, ~~2022~~

Core Java - Concepts, multithreading, collections, design patterns, JDBC, SpringBoot

J2EE - JSP objects, Servlet lifecycle, framework MVC

Rest - Go through interview questions.

Design Patterns - Singleton, Factory pattern,

MVC - Struts & Hibernate

HDFS - Big data concepts, HDFS itself

file read/write & configurations

hdfs vs old style

Storage policy, replication factor, chunk size

Mapreduce - functioning

configuration

parallel processing

interview question

Yarn - Mapreduce vs YARN

Apache Hadoop, components of hadoop, & hadoop2

Hive - Architecture, modes, Tez vs MapR

hive vs SQL / Traditional RDBMS

compression  
SNAPPY

hive vs MapReduce

partitioned file format  
hive vs Pig

hive queries, UDAF Evaluate method

Map join, self join, cross join,  
UDF, UDAF, aggr. functn  
purpose ← validity

Hcatalog, complex types, Serde

Pig - Architecture

modes, functions, complex types, UDF

extend EvalFunc  
exec method

$*/10 \quad 11-13 \quad * \quad * \quad *$   $\Rightarrow$  Every 10 min bet<sup>n</sup> 11 to 13.59 pm

Cronjob -  $0 \quad * \quad * \quad * \quad *$   $\Rightarrow$  0th minute every hour.

→ minute (0 - 59)

→ Hour (0 - 23)

→ Day of month (1 - 31)

→ Month (1 - 12)

→ day of week (0 - 6)  $\nearrow$  sunday = 0 or 7

⇒ \* \* \* \* \* command.

## Spark SQL

- optimization possible
- control over structure of data
- files, JDBC, Hive, Parquet
- Dataset<Row>  $\Rightarrow$  Dataframes  $\rightarrow$  limitation
- Encoder  $\Rightarrow$  works like SerDe
  - Generate a bytecode which can interact with off-heap data.
  - no need to deserialize whole object
- DataSet  $\Rightarrow$  compile time safety.

Features  $\Rightarrow$  ① Optimization

② Compile time check

③ Persistence

④ Convertible  $\Rightarrow$  toDF(), toDS()

⑤ fast ⑥ less memory

⑦ Sample API for Java & Scala

- RDD vs DF vs DS

$\Rightarrow$  RDD  $\Rightarrow$  RDD.toDF  $\Rightarrow$  define case class

gives DP

or Dataset<Row>

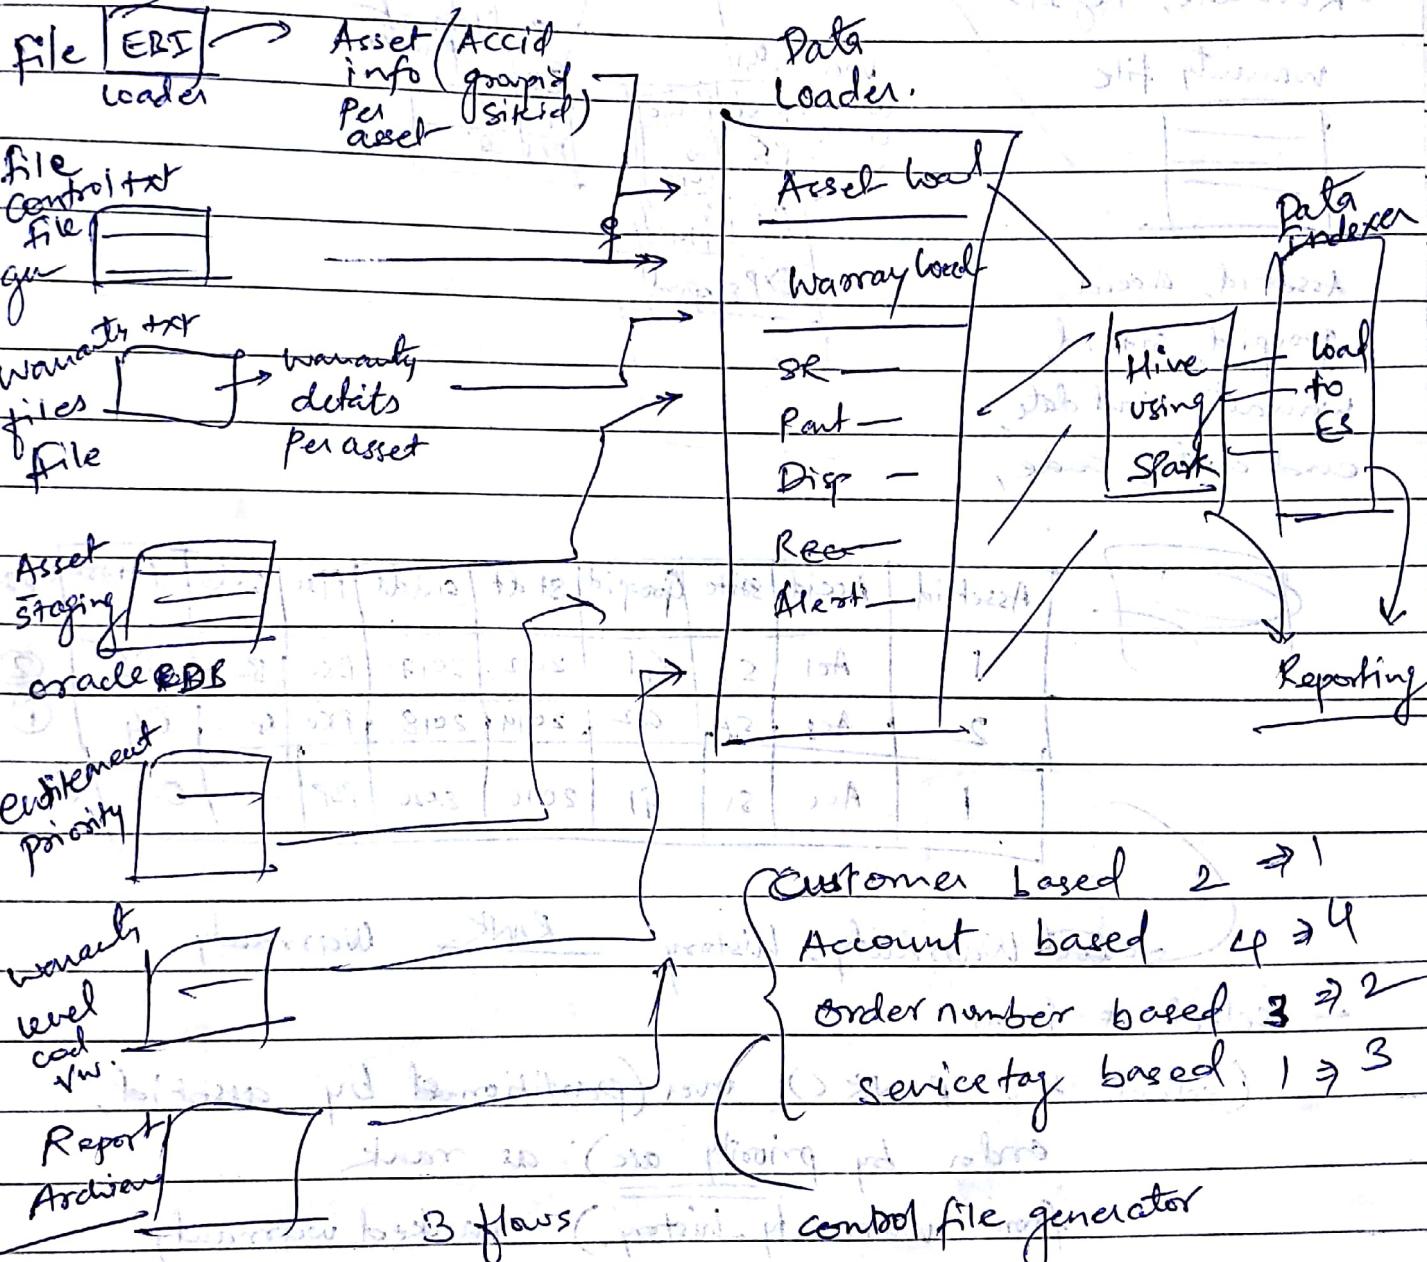
DF.as [Employee]

gives dataset<Employee>

## Dell PPRC.

- Dell leading computer peripherals manufacturer
- wanted a way to organize their repository in hadoop cluster
- 7 components → Asset, warranty, Service Request, Part, dispatch, recommendation, alerts

8 node  
24 TB  
dat



1 → enrichment

Pending

2 → modified enriched

3 → moved or deleted

his history  
to check what was here data for asset warranty

ESF file



Spark

-Join the dataset

Check for validity

-join with oracle

Create dataset to load into

hive.

& Create RANK over it.

Assetid, Accid,  
groupid, siteid,  
start date, end date  
SR LVL code, Reg ID

warranty file



warranty

bit  
code  
ver.

BSC	Basic
PS	Pno
	SUPP
PSP	Pro
	plus
	Exp & Expired

entitlement

BSC	5
PS	4
PSP	3

Assetid, Accid,  
groupid, siteid,  
warranty start date,  
end date, code,

Assetid	Accid	site	Groupid	st_dt	end_dt	Plan	Code	Priority	RANK
1	Acc1	s1	G1	2012	2017	BSC	BSC001	05	②
2	Acc2	s2	A2	2014	2018	PKO	PKO...	04	①
1	Acc1	s1	G1	2016	2020	PSP	---	03	①

(Total number)

Get Warranty\_History  $\rightarrow$  RANK  $\rightarrow$  Warranty.

Select \* from

(Select \*, RANK()) over (partitioned by assetid,  
order by priority asc) as rank

Rank from (Warranty\_History) as ranked\_warranty

where ranked\_warranty.rank = 1;

$\Rightarrow$

Priority asc  $\rightarrow$  Priority Asc, Start DESC, End DT DESC.

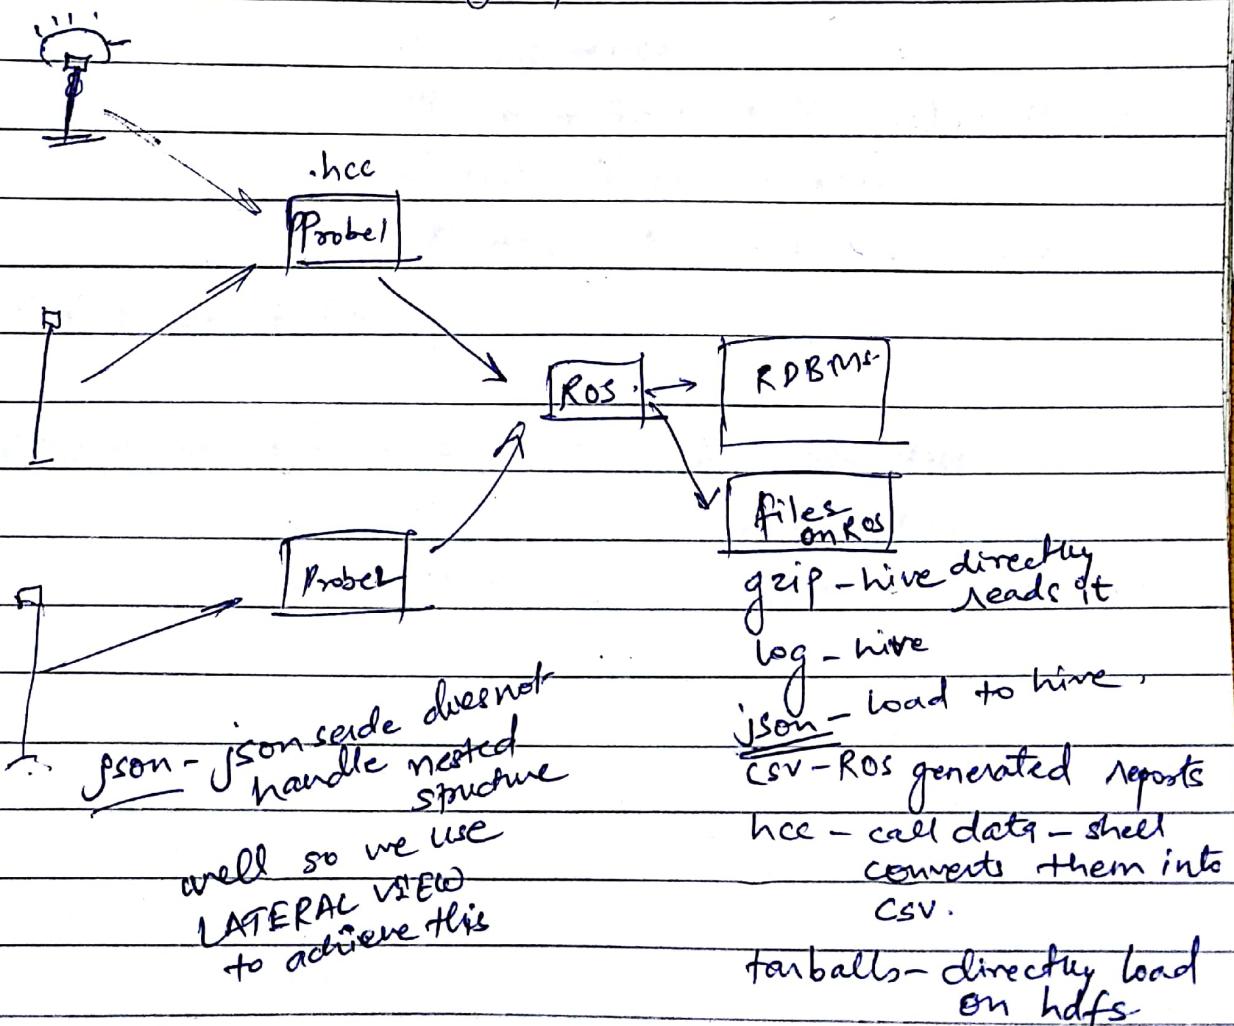
Priority high last.

# Empirix Intellisight & Holistix

Empirix provides carrier grade products & services to major network providers. Their clients were Vodafone, O2, Tuk. So this project was to design an initial migration project of the application from traditional RDBMS to hadoop infra.

- Tasks :- identified what all sources for data.
  - what all types of input data.
- Steps :- shell scripts to read them & check its validity → HDFS
  - Hive scripts to load them into hive & process
  - RDBMS → sqoop scripts → Hive tables
  - Java module to create ladder diagram out of this data.
  - cron job → nightly

6 nodes  
5TB

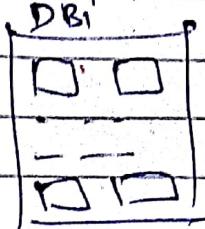


grid  
wire  
policy  
center  
Grid desire

S4X  
Ingestion

Denormalized  
form DATE: [ ] [ ] [ ] [ ]

S4  
Curation



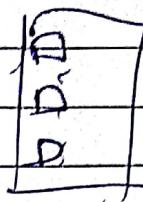
gwpc



gwpc

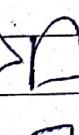
denormalized  
+ lookup.

target



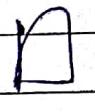
gwcc

+  
lookup

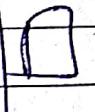


dedup the  
data

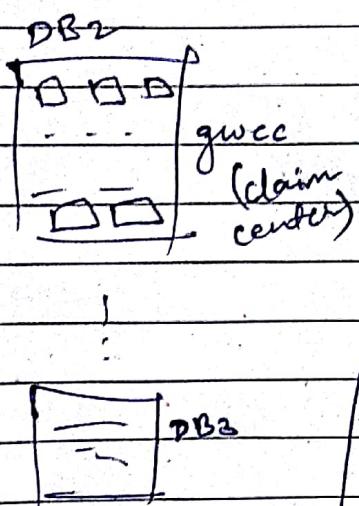
+  
ETL



+  
Auditing



+  
Historical  
& Ince-  
mental  
load

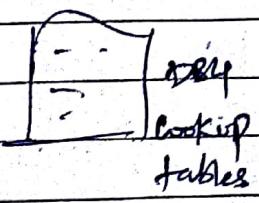


gwcc  
(claim  
center)

DFD  
framework.  
Ingestion  
into  
hadoop

Batchid

Only Append



only  
lookup  
tables

## Challenges

- ① State auto - orc file reading from spark vs reading from hive  
- spark printed column values as Null due to  
orc version was older  
Action - we installed newer orc version packages to spark  
Solution libraries.
- ② Spark Jobs was taking too long,  
Solution - Performance enhancements using broadcast join,  
better spark config. & StreamTable hints in sparkSQL  
Groupby key removed by reduceByKey
- ③ OOM errors - removed collects wherever possible  
- Also used mapPartitionsWithIndex & then  
collected at partition level & saved data
- ④ Spark version mismatch in dev vs UAT environment.  
- this caused the job to fail due to difference in  
lineage (DAG) created by job  
- Made UAT cluster on Spark version as Dev
- ⑤ OOM - due to improper join sequences (Largest table last)
- ⑥ Architecture challenge - Most cost on cluster. but job ETL  
time was 6 hours-  
Solution - Suggested dataproc cluster transient one  
& a small permanent hive dataproc cluster for hive
- ⑦ Difficult to debug the code- Solution - we added isDebugMode  
flag & printed 20 records for better debugging.

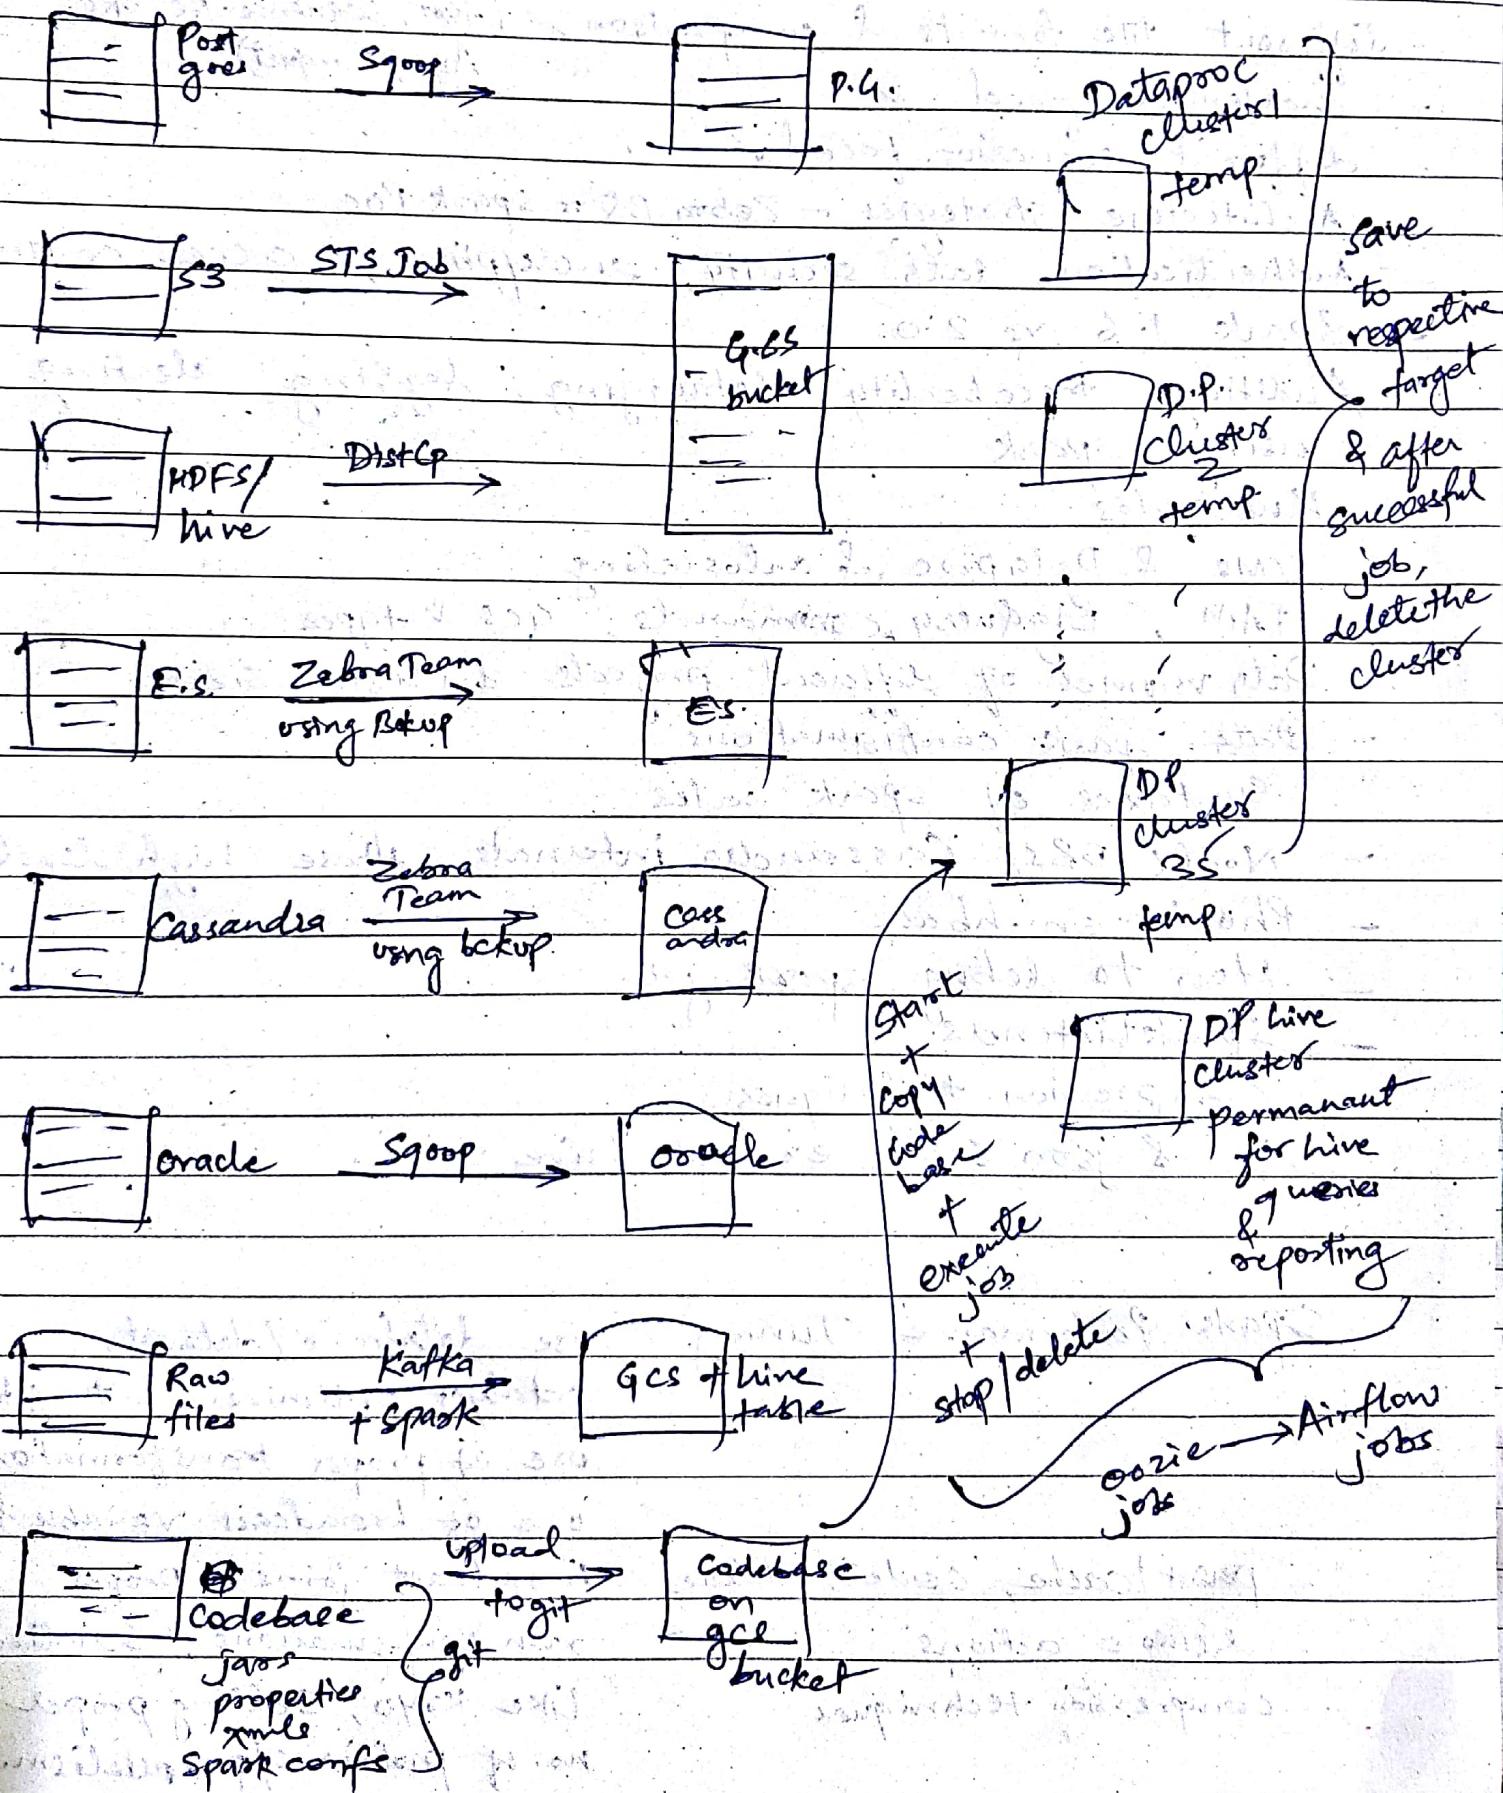
DATE:

- different file formats & comparison ✓ (text, sequence, RC, ORC, Avro, Parquet)
- challenges faced ✓
- different scenarios faced ✓
- Architecture strategies - Zebra BQ vs spark Poc
- Authentication, data security, encryption, access control
- Spark 1.6 vs 2.0
- Auditing, traceability, debugging, logging, alerting using spark
- Kubernetes
- VMs & Data proc & autoscaling
- IAM, Bigquery commands, GCS & types
- Data volumes of different projects & cluster sizes.
- ~~Data~~ Spark configurations
- Syntaxes of spark code
- NoSQL DBs, Cassandra internals, Hbase high level
- Phoenix on hbase
- How to debug spark job.
- SparkListenders
- Compression techniques
- XML & json serde example hive

Spark Performance Tuning - use dataframe/dataset, catalyst optimizer & tungsten, use of proper transformation, use of broadcast variable & broadcast joins, proper persist/cache, code modularize, remove actions, compression techniques, like Kryo, setting proper no. of partitions/parallelism, run configurations, Serialize

## Zebra Migration

DATE: [ ]



## Zebra ReArch vs DA

DataLake

mdm DB

Mdm tables

EDS DB

exposed view

mdm views

EDS views

BQ script

to fetch  
data from

mdm &  
EDS views

for  
snapshot  
& historical  
events

DATE:

□ □ □ □

Comparison

spark (BQ)

t-gms-cdm

Aggregated  
CDM table

BQ

script

BQ

Hadoop vs Dataproc :- Support by google, one place of everything, can be shutdown, GS + dataproc combination

Dataflow vs Spark

- ① Resource mgmt - Dataflow allocates on demand
- ② Interactive shell - Spark has this
- ③ Streaming feature - dataflow better
- ④ Complexity of coding - dataflow for simpler transforms

GS, Dataproc, BigQuery, Cloud SQL, Cloud Spanner,  
Cloud BigTable, Cloud Datastore, Dataflow, Dataprep,  
Stackdriver, Kubernetes, Compute Engine, App Engine,  
Pub/Sub, Interconnect, STS, Load balancing,  
Cloud Functions (trigger based), Google Cloud Storage,  
types & Billing, BQ Billing,